Workshop: LLVM for Reverse Engineers





Duncan Ogilvie

Setup: workshop.ogilvie.pl

- 1. Login to your GitHub account
- 2. Fork the repository
- 3. Click the green <> Code button
- 4. Press ... and then New with options...
- 5. Change Machine type to 4-core
- 6. Then Create codespace
- 7. Wait ~3 minutes while the image is loading 🥌
 - Press Show log to see progress

Introduction

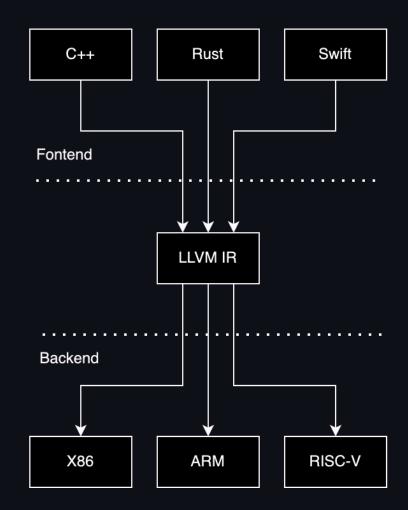
- Meant for absolute (LLVM) beginners
 - Basic familiarity with C and assembly useful
- Format: hands-on workshop
- Interactive
- Available for on-site training: training@ogilvie.pl

Who am I?

- Creator of x64dbg
- Worked in DRM for 5 years
- Currently working in mobile security R&D

What is LLVM IR?

- Low Level Virtual Machine (paper)
 - Authors: Chris Lattner, Vikram Adve (2002)
 - Meant for compiler development
- Intermediate Representation (IR)
 - Platform agnostic (mostly)
 - Functions, basic blocks, control flow, etc.
 - Reduced Instruction Set Computer ('RISC')
 - Single Static Assignment (SSA)
- Reusable optimization/code generation pipeline



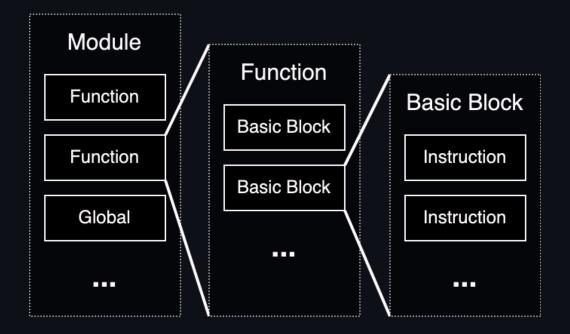
Why LLVM IR?

- Lifting
 - Deobfuscation
 - Emulation
 - Binary rewriting
- Fuzzing/analysis (vulnerability research)
- Obfuscation

Note: tame your expectations.

LLVM IR: Module

- Functions
 - Basic Blocks
 - Instructions (reference)
- Globals (variables)
- Metadata



LLVM IR: Hello World

C:

```
int hello(int x, int y) {
  return x + y;
}
```

LLVM IR:

```
define i32 @hello(i32 %0, i32 %1) {
2: ; entry block label (optional)
    %3 = add i32 %0, %1
    ret i32 %3
}
```

- Identifiers: @global, %local
- No signed/unsigned number types
- Implicit numbering vs explicit naming of values

LLVM IR: Clang

hello.c:

```
define dso_local i32 @hello(i32 noundef %0, i32 noundef %1) #0 {
 %3 = alloca i32, align 4
 %4 = alloca i32, align 4
  store i32 %0, ptr %3, align 4
  store i32 %1, ptr %4, align 4
 %5 = load i32, ptr %3, align 4
 %6 = load i32, ptr %4, align 4
 %7 = add i32 %5, %6
  ret i32 %7
attributes #0 = { noinline nounwind optnone uwtable ... }
```

- dso_local: Runtime Preemption Specifier (always emitted by Clang)
- noundef: Parameter attribute to indicate the value is always defined
- attributes: Group of function attributes

LLVM IR: Single Static Assignment (SSA)

- Local values are defined, not assigned
 - Variable is a misnomer
 - Also called registers
- You cannot define the same value twice
- Does not apply to memory
 - The load and store instructions operate on ptr values

LLVM IR: Control Flow

cfg_alloca.ll:

```
uint32_t cfg(uint32_t x) {
  if (x > 10) return 123;
  else     return 321;
}
```

LLVM IR: phi

```
cfg_phi.ll:
```

```
define i32 @cfg(i32 %x) {
 %cond = icmp ugt i32 %x, 10
 br i1 %cond, label %bb_if, label %bb_else
bb_if: ; preds = %entry, x > 10
 %result_if = add <u>i32</u> 0, <u>123</u>
 br label %bb end
bb_else: ; preds = %entry, !(x > 10)
 %result_else = add i32 0, 321
 br label %bb end
%result = phi i32 [%result_if, %bb_if], [%result_else, %bb_else]
 ret i32 %result
```

Simplifies analysis/optimization passes.

LLVM IR: select (ternary)

```
cfg_select.ll:
```

```
define i32 @cfg(i32 %x) {
   %cond = icmp ugt i32 %x, 10
   %result = select i1 %cond, i32 123, i32 321
   ret i32 %result
}
```

```
uint32_t cfg(uint32_t x) {
  return (x > 10) // cond
    ? 123
    : 321
}
```

LLVM IR: Optimizations

- Compilers emit non-optimal LLVM IR
 - Often easier to write a compiler this way
- LLVM IR was designed to be suitable for optimizations
- The compiler has a pipeline of *optimization passes*
 - Every pass has its' own purpose

LLVM IR: Exercises (part 1)

Instructions: exercises/1_llvmir/README.md (Exercise 1a-1d)

LLVM IR: getelementptr

- Purpose: pointer arithmetic
 - Arrays
 - Structs
 - Used absolutely everywhere
- Does not read memory
- Most confusing instruction
 - Similar to the x86 lea instruction

LLVM IR: getelementptr (array)

C:

```
uint32_t arrayExample(uint32_t* arr) {
  return arr[5];
}
```

LLVM IR:

```
define i32 @arrayExample(ptr %arr) #0 {
   %ptr_idx_5 = getelementptr i32, ptr %arr, i64 5
   %result = load i32, ptr %ptr_idx_5
   ret i32 %result
}
```

• $ptr_idx_5 = (uintptr_t)arr + 5 * sizeof(i32)$

LLVM IR: getelementptr (member)

C:

```
typedef struct { uint64_t a[2]; uint32_t b; uint32_t c[5]; } MyStruct;
uint32_t structExample1(MyStruct* s) {
  return s->b; // s[0] b
}
```

LLVM IR:

```
%struct.MyStruct = type { [2 x i64], i32, [5 x i32] }

define i32 @structExample1(ptr %s) #0 {
   %ptr_b = getelementptr %struct.MyStruct, ptr %s, i32 0, i32 1
   %result = load i32, ptr %ptr_b
   ret i32 %result
}
```

• ptr_b = (uintptr_t)s + 0 * sizeof(MyStruct) + offsetof(MyStruct, b)

LLVM IR: getelementptr (member array)

C:

```
typedef struct { uint64_t a[2]; uint32_t b; uint32_t c[5]; } MyStruct;
uint32_t structExample2(MyStruct* s) {
  return s->c[3];
}
```

LLVM IR:

```
%struct.MyStruct = type { [2 x i64], i32, [5 x i32] }

define i32 @structExample2(ptr %s) #0 {
   %ptr_c = getelementptr %struct.MyStruct, ptr %s, i32 0, i32 2
   %ptr_c_3 = getelementptr [5 x i32], ptr %ptr_c, i32 0, i32 3
   %result = load i32, ptr %ptr_c_3
   ret i32 %result
}
```

LLVM IR: getelementptr (optimization)

```
%struct.MyStruct = type { [2 x i64], i32, [5 x i32] }
define i32 @structExample2(ptr %s) #0 {
  %ptr_c = getelementptr %struct.MyStruct, ptr %s, i32 0, i32 2
  ptr_c_3 = getelementptr [5 x i32], ptr <math>ptr_c, i32 0, i32 3
  %result = load i32, ptr %ptr_c_3
  ret i32 %result
define i32 @structExample2_opt(ptr %s) #0 {
  %ptr_c_3 = getelementptr %struct.MyStruct, ptr %s, i32 0, i32 2, i32 3
  %result = load i32, ptr %ptr_c_3
  ret i32 %result
```

LLVM IR: getelementptr (flattening)

```
%struct.MyStruct = type { [2 x i64], i32, [5 x i32] }
define i64 @structExample3(ptr %s) #0 {
  %ptr_a = getelementptr %struct.MyStruct, ptr %s, i32 0, i32 0
  ptr_a_1 = getelementptr [2 x i64], ptr <math>ptr_a, i32 0, i32 1
  %result = load i64, ptr %ptr_a_1
  ret i64 %result
define i64 @structExample3_opt(ptr %s) #0 {
  ; No reference to MyStruct at all anymore
  ptr_a_1 = getelementptr [2 \times i64], ptr %s, i64 0, i64 1
  %result = load i64, ptr %ptr_a_1
  ret i64 %result
```

- LLVM IR Godbolt
- C Godbolt (play with the optimization settings)

LLVM IR: Exercises (part 2)

Instructions: exercises/1_llvmir/README.md (Exercise 2a)

Quick break (15 min)





What is Remill?

- Authors: Trail of Bits (2015)
- Lifts native instructions to LLVM IR
 - Applications: binary analysis/instrumentation/emulation
 - o Architectures: ARM, X86, PPC, SPARC, Sleigh
- Mild abuse of the IR, requires some tricks

Remill: Concepts

- Instruction semantics in C++
 - Easier to maintain
 - Compiled to LLVM IR
- State* structure -> CPU Registers
- Memory* pointer -> memory manager
 - Total ordering to preserve semantics
- 'Massaging' required

Remill: Instruction Semantics

Semantics of the x86 mov instruction:

```
template <typename D, typename S>
DEF_SEM(MOV, D dst, const S src) {
   WriteZExt(dst, Read(src));
   return memory;
}

DEF_ISEL(MOV_GPRv_MEMv_32) = MOV<R32W, M32>;
```



Remill: Lifting Basic Blocks

Basic Block Definition:

```
Memory *__remill_basic_block(State &state, Ptr block_addr, Memory *memory);
```

- Calls to the semantics are inserted here.
- State is fully symbolic
- Requires additional work to restore the calling convention

Remill: High level example

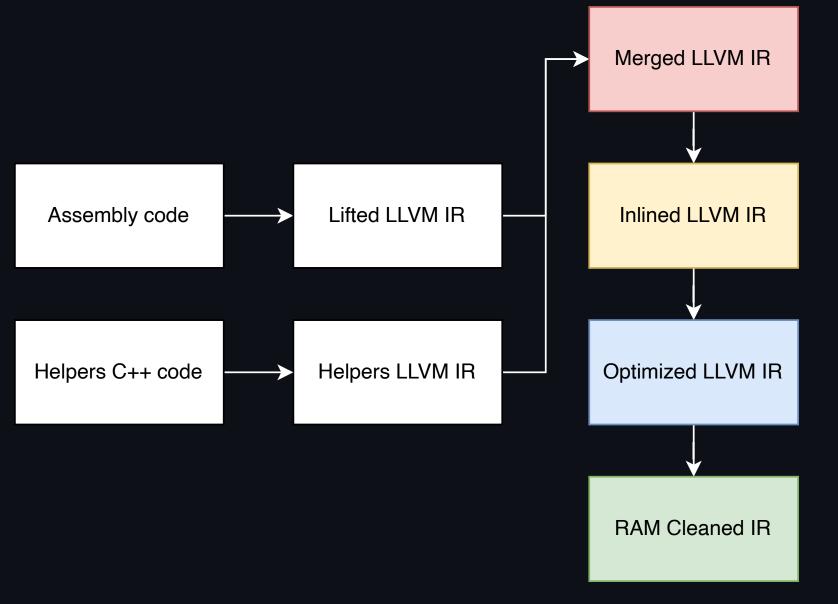
```
Memory *__remill_basic_block(State &state, Ptr block_addr, Memory *memory) {
    // mov rax, rdi
    state.rax = state.rdi;
    state.rip += 3;
    // ret
    state.rip = *(Ptr*)state.rsp;
    state.rsp += sizeof(Ptr);
    return __remill_function_return(state, state.rip, memory);
}
```

Remill: Helpers

```
Memory *__remill_write_memory_8(Memory *m, addr_t a, uint8_t v);
Memory *__remill_write_memory_16(Memory *m, addr_t a, uint16_t v);
Memory *__remill_write_memory_32(Memory *m, addr_t a, uint32_t v);
Memory *__remill_write_memory_64(Memory *m, addr_t a, uint64_t v);
```

- Abstraction to represent interaction with the host CPU (memory, calls, indirect branches, syscalls, flag computations)
- Implementation varies depending on the purpose (emulation, symbolic execution, decompilation)
- Makes the lifted IR difficult to work with for humans (extremely verbose)

Lifting Overview



Remill: Exercises

Helpers: helpers/x86_64/RemillHelpers.cpp

Instructions: exercises/3_lifting/README.md

Closing Remarks

- Continue at home!
- Thanks: Matteo Favaro
- Get in touch: training@ogilvie.pl

Bonus: links

- KLEE Symbolic Execution Engine
- SymCC
- SATURN
- Tickling VMProtect with LLVM
- VMProtect devirtualization using Triton/LLVM
- souper
- RetDec
- Tigress transformations
- Modern Obfuscation Techniques
- RISC-Y Business: Raging against the reduced machine